Contents

Chapter 1

Software Process for Multiphysics Multicomponent Codes

1.1 Introduction

The computational science and engineering (CSE) communities have a
mixed record of using software engineering and adopting good software prac-
tices. Many codes adopt software practices when size and composition of the
code stops progress with adopting them. In rarer instances code projects start
with an awareness of the importance of software process and build it in from
the beginning. As more codes have crossed the threshold of being manageable
without software engineering they have increasingly been adopting software
processes derived from outside the scientific domain. The driving force behind
adoption is usually the realization that without using software engineering
practices, the development, verification and maintenance of code becomes in-
tractable. State-of-the-art for software engineering practices in CSE codes lags
behind that in the commercial software space. There are many reasons for it,
but an important one is that not all software engineering practices can be
adopted by the CSE code developers. And there is very little research in the
software engineering community targetted to specific needs of CSE codes.

many software best practices are not well-suited for CSE codes without
modification and/or customization, in particular to multiphysics multicom-
ponents codes that run on the large HPC platforms. Sometimes the inherent
physics of scientific applications require different software methodologies, at
others a premium is placed on performance rather than code architecture.
Still others are more sociological and due to the type of institutions where
such codes are developed. The challenges for scientific applications range from
their architecture to the process for their maintenance and growth. The stan-
dard practices adopted by the CSE codes include repositories for code version
control, modular code design, licensing process, regular testing, documenta-
tion, release and distribution policies and contribution policies. Less frequently
used practices include code-review and code-depracation. Agile developmenet
methods tend to have limited usefulness in the CSE lifecycle, only a select
subset of those have met with success. The degree of adoption and sophistica-
tion in using these practices varies among teams. Many of the reasons behind

less penetration are discussed in section **??**. Even among the widely adopted practices, most are modified and customized by the developers for their own needs. The next two sections outline the challenges that are either unique to, or are more dominant in this domain than elsewhere.

1.2   Lifecycle

Scientific software is designed to model some phenomena in the physical world. The phenomena may be at microscopic level, for example protein folding, or at really large scales, for example galaxy cluster mergers, or have a range of scales. We use the term physical to also include chemical and biological systems since physical processes are underlying building blocks for those systems too. Their physical characteristics are translated into mathematical models that are said to describe the essential features of the behaviour of the system being studied. These equations are then discretized, and numerical algorithms are used to solve them. In general there are many more degrees of freedom in the development and lifecycle of scientific software that are not encountered elsewhere.

1.2.1   Development Cycle

For scientific simualtions, modeling begins with equations that describe the general class of behavior to be studied, for example the Navier-Stokes equations describe the flow of compressible and incompressible fluids, and Van-der-vaal equations describe force interactions among molecules in a material. There may be more than one set of equations if there are behaviors that are not adequately captured by one set. In translating the model from mathematical representation to computational representation two processes go on simultaneously, discretization and approximation. One can argue that discretization is by definition an approximation because it is in effect sampling continuous behavior where information is lost between sampling intervals. This loss shows up as error terms in the discretized equations. But error terms are not the only approximations. Depending upon the level of understanding of specific sub-phenomena, and available compute resources, scientists also use their judgement to make other approximations. Sometimes, to focus on a particular behavior, a term in an equation may be simplified or may be even completely dropped. At other times some physical details may be dropped from the model because they are not understood well enough by the scientists.

The next stage in developing the code is finding the appropriate numerical methods for each of the models. Sometimes good methods exists that can be used " as-is", at others they may need to be customized, or new methods
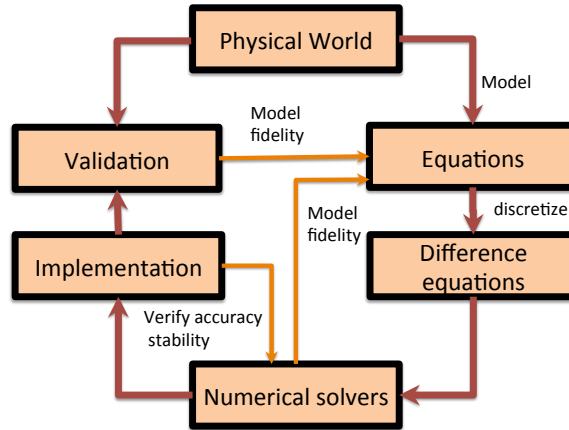
FIGURE 1.1: Development cycle of modeling with partial differential equations

may need to be developed. There may need to be validation of the method's applicability to the model if the method is new or significantly modified. Unless an implementation of the method is readily available as a third party software (stand-alone or in a library), it has to be implemented and verified. It is at this stage that the development of a CSE code begins to resemble that of general software. The numerical algorithms is specified, the semantics are understood, and they need to be translated into executable code. Figure **??** gives an example of the development cycle of a multiphysics application modeled using partial differential equations.

1.2.2 Verification and Validation

There are many stages in the development cycle of scientific software where errors can be introduced. Classes of these errors are mostly orthogonal to one another, a good verification and validation methodology will reflect and exploit that. The terms verification and validation are used in ways where they don't always mean the same thing. In one narrow definition of validation, it ensures that the mathematical model correctly defines the physical phenomena, while verification makes sure that the implementation of the model is correct. In other words, model is validated against observations or experiments from the physical world, whereas verification encompasses all other forms of testing. Other definitions give broader scope to validation. For example, validation can also apply to a numerical method through exercises such as code-to-code comparisons, and its order can be validated through convergence studies. Sim-

ilarly the implementation of a solver can be validated against an analytically obtained solution for some model if the same solver can be applied and the analytical solution is also known. Though this is not always possible. Irrespective of specific defintions, what is true is that correctness must be assured at all the stages from model to implementation.

There are many degrees of freedom in the process of deriving a model as discussed in the previous section, therefore, the validation of the model must be carefully calibrated by scientific experts. Similarly, verification of a numerical method requires applied math expertise because the method needs verification of its stability, accuracy and order of convergence in addition to correctness. Numerical methods have their own error analysis because of approximations. Many of these methods are themselves objects of ongoing research, so their implementation may need modifications from time to time. Whenever that happens the entire gamut of verification and validation needs to be applied again. This is an instance of particular challenges in the CSE software where no amount of specification is enough to hand the implementation over to the software engineers/developers. A close collaboration is a must because the process has to be iterative with scientific judgement applied at every iteration.

One other unique verification challenge in CSE software is the consequences of finite machine precision of floating point numbers. Any change in compilers, optimization levels, and even order of operations can cause numerical drift in the solutions. Especially in applications that have a large range of scales it can be extremely difficult to differentiate between a legitimate bug and a numerical drift. Therefore relying upon bitwise reproducibity of the solution is rarely a sufficient method for verifying the continued correctness of an application behavior. Robust diagnostics (such as statistics or conservation of physical quantities) need to be built into the verification process. Knowing the error bars also helps. This issue is discussed in greater detail in chapter **??**.

1.2.3   Maintenance and Extensions

A typical software project undergoes a design and development phase, followed by production and maintenance phase. This clearly defined lifecyle rarely applies to scientific codes. In most successful multiphysics codes the infrastructure and solver API's are the only entities that have a distinct development phase which does not spill over into the remainder of the lifecycle. The development of CSE software is usually responding to an immediate scientific need, so the codes get employed in production as soon as a minimal set of computational modules necessary for even one scientific project are built. Similarly, the development of computational modules almost never stops all through the code lifecycle because new findings in science and math almost continuously place new demands on them. The additions are mostly incremental when they incorporate new findings into an existing feature, they can

also be substantial when new capabilities are added. The need for new capabilities may arise from greater model fidelity, or from trying to simulate a more complex model. Sometimes a code designed for one scientific field may have enough in common with another field that capabilities may be added to enable it for the new field.

Whatever may be the cause, co-existence of development and production/maintenance phases is a constant challenge to the code teams. It becomes acute when the code needs to undergo major version changes. The former can be managed with some repository discipline in the team coupled with a solid testing regime. The latter is a much bigger challenge where the plan has to concern itself with questions such as how much backward compatibility is suitable, how much code can go offline, and how to reconcile ongoing development in code sections that are substantially different between versions. FLASH's example in section **??** describes a couple of strategies that met the conflicting needs of developers and production users in both scenarios. Both required co-operation and buy-in from all the stakeholders to be successful.

1.2.4   Using CSE Software

There is a fundamental requirement from the users of scientific software that rarely comes into play for users of other kinds of software. For good results, the users of scientific software cannot treat it as a black box. They must understand the models well. They must also know and understand the range of applicability of numerical algorithms to their physical regimes, and also the accuracy and stability behavior of the algorithms. It is very possible to apply the methods in inappropriate ways and obtain scientifically useless results. Even worse, one may obtain wrong results and not even know that they are wrong. This is because some phenomena are very sensitive to perturbations. If one applies a method without sufficient resolution, the perturbations may be filtered out and the outcome of the simulation may be physically valid while being completely wrong for the phenomenon being studied. Similarly, sometimes equations have mathematically valid but physically invalid solution. A badly applied numerical scheme may converge to such a solution. Even though in this situation it becomes obvious that the solution is not right, it may happen after significant wasteful use of resources. These are some of the reasons that also play a role in tendency of scientific codes to do strict gatekeeping for contributions, and mostly operate in the cathedral mode.

1.3   Domain Challenges

Multiphysics codes, by their definition, have more than one mathematical model that they are solving. A typical code combines 3-4 diverse models, the

more extreme ones may employ as many as a dozen. In a rare calculation all models work with the same discretization using similar algorithmic approach (for instance stencil computations in explicit PDE solvers). More common is to have models with diverse discretizations and algorithms. Each operator has its own preferred data layout and movement, and it usually differs from those needed by the other operators. Normally these challenges can be mitigated through encapsulation and well defined API's. The outer wrapper layers of the operators can carry out data transformations as needed. There are two factors against taking this approach in CSE codes: (1) physics is not always friendly to encapsulation, and (2) the codes are performance sensitive and wholesale data movement significantly degrades performance.

The CSE simulation codes model the physical world which does not have neat modularization. Various phenomena have tightly coupled dependencies that are hard to break. These dependencies and tight couplings also translate into their mathematical models and it becomes hard to eliminate lateral interactions among code modules implementing the models. An attempt to force encapsulations by hoisting up the lateral interactions to the API level can explode the size of the API. And if not done carefully this can also lead to extra data movement. The module designs, therefore, have to be cognizant of potential lateral interactions and make allowances for them. Similarly the data structures have to take into account the diverse demands placed on them by different operators and carefully consider the trade-offs during software design. Considerations such as these are not common in software outside of CSE.

In a CSE software design, separation of concerns is of utmost importance. Orthogonalization of expertise requirements into different code components allows developers to focus on what they know best. Another natural fallout of this approach is that different dimensions of complexities in the algorithm space are handled separately. The numerical algorithms associated with physics operators are complex because of accuracy and stability concerns, and require mathematical expertise. They are not logically as complex. Whereas machinery for managing the discretizations and interoperability among code components is likely to be less complex numerically, but could be very complex logically. A third axis of concern is parallelization, which brings in some features that are unique to CSE codes, such as domain decomposition, aspects of synchronization and dependencies, and performance impact of the design choices. With appropriate separation of concerns not only do these aspects of software development not interfere with one another, they help make the development tractable.

Multiphysics multiscale codes are almost unique in the software world in their need to tightly integrate third party software, which comes in the form of numerical libraries. Because multiphysics codes combine expertise from many domains, the numerical solvers they use also require diverse applied mathematics expertise. It is nearly impossible for any one team to assemble all the necessary expertise. The only form in which such expertise is encoded is in the form of numerical libraries or other third party software. The developers are

faced with a choice between integrating third party software, or developing their own technology which is likely to yield inferior quality of solution. However, as mentioned in section **??**, the use of third party does not absolve them from understanding its appropriate use. Additionally, information about appropriate use of third party software within the context of a larger code must also be communicated to the users of the code.

Testing of CSE software needs to reflect the layered complexity that the codes themselves have. The first line of attack is the unit tests where possible. However, as mentioned in chapter **??**, some dependencies are meaningless to break in mathematical software, and where such dependences exist a unit test cannot be devised. Testing relies upon no-change-within-bounds tests where minimum possible combination of units are used within the dependence constraints. In effect these minimally combined tests play the same role in the testing regime that unit tests do because they focus on possible defects in a very narrow section of the code. Multicomponent CSE software have to test various permutions and combination of components in many different ways. These tests verify that all configurations are within the accuracy and stability constraints.

Another unique aspect of multiphysics CSE software is its need for performance portability. HPC machines are expensive and rare resources, they need to be used efficiently. In order to do that codes should ideally be optimized for each machine. However, typical lifecycle of a multiphysics software spans many generations of HPC platform lifecycle, which is about 3-4 years. Depending upon the size of the code, optimization for a specific target platform can take a significant fraction of the platform lifecycle. During the optimization phase, the code is not available for science. Thus a large fraction of scientists' time is lost in porting and optimizing the code over and over. Another dimension of this problem is that even within the same generation the platforms differ from one another. So machine specific optimization ties a code to one machine. These factors make platform-specific optimizations unattractive. Instead, HPC CSE codes consider the trade-offs and opt to design their software using constructs that perform modestly well across a range of platforms.

1.4   Institutional Challenges

Many adaptations in the software engineering described in the previous section pertained to software design and testing. In particular they spoke to challenges of modularity, performance and unit-testing because of the intertwined nature of the problems being tackled by these codes. However, all challenges faced by the CSE codes are not because of the nature of problems they solve. Many challenges are specific to the kind of organizations and the research communities where these codes are developed. The most

crippling and pervasive challenge faced by CSE codes in general, and multiphysics codes in particular, is that they rarely get funding for development of software infrastructure. There is evidence that when software is designed well it pays huge dividends in scientific productivity from the miniscule number of projects that secured such funding for software infrastructure design. Even with the evidence the scientific establishment remains unconvinced about the criticality of investment in software engineering. The funding is carved out of scientific goal oriented projects that have their own priorities and time-line. This model often ends up short-changing the software engineering.

The scientific output is measured in terms of publications which in turn depend upon the data produced by the simulations. Therefore in a project driven purely by scientific objectives, the short-term science goals can lead to compromise on the quality of software design. Quick-and-dirty often triumphs over long term planning. The cost of future lost productivity is not appreciated until it is too late. By the time design deficiencies are realized usually the software has grown too large to remove the deficiencies in any easy way. Software engineering is forcibly imposed on the code, which it at best a band-aid solution. This is another reason why many of the software practices are not embraced by the CSE community.

Another instituitional challenge faced by SE for CSE is the training. Multiphysics codes require a broad range of expertise in domain science from their developers, software engineering is an added requirement. The developers are not trained in SE, many learn them on the job through reading or talking to colleagues. The practices are applied as they understand them, usually picking only what is of most importance for their own development. This can be both good and bad. Good because it sifts out the unnecessary aspects of the SE practice, and bad because it is not always true that the sifted out aspects were really not necessary. It just means that the person adopting the practice did not understand how to use them, or their importance.

Institutional challenges also arise from scarcity of resources and stability. The domain and numerical algorithmic expertise is rarely replicated in a team developing the multiphysics CSE application. Even otherwise, deep expertise in the domain may be needed to model the phenomenon right, and that kind of expertise is hard to come by. Then there is the challenge to communicating the model to the software engineer if there is one on the team, or to other members of the team with some other domain expertise. It requires at least a few developers in the team who can act as interpreters for various domain expertise and are able to integrate them. Such abilities take a great deal of time and effort to develop, neither of which are possible in the academic institutions where these codes are typically organically grown. The available human resources in these institutions are post-docs and students who move on, so there is no retention of institutional knowledge about the code. A few projects that do see the need for software professionals struggle to find ways of funding them or to provide a path to professional growth.

The above institutional challenges also provide a clue about why any set

software development methodology is hard, and often even undesirable, to adopt in such projects. For example, the principles behind the agile manifesto apply, but not all the formalized process does. These codes are developed by interdisciplinary teams where interactions and collaborations are preferred over process. The code development and its use for science go on in parallel, so the requirements change and there is quick feedback when they do. For the same reason, the code needs to be in working condition almost all the time. However, scarcity of resources does not allow the professional roles in the agile process to be played out efficiently. There is no clear separation between the developer and the client, many developers of the code are also the scientists who use it for their research. Because software development goes hand in hand with research and exploration of the algorithms it is impossible to do either within fixed time-frames. This constraint effectively eliminates using sprints. The waterfall model is even less useful because it is impossible to do a full specification ahead of time. The code has to grow and alter organically as the scientific understanding grows, the effect of using technologies are digested and requirements change.

The need for deep expertise, and the fact that the developer of a complex physics module is almost definitely going to leave with possibly no replacement, documentation of various kind takes on a crucial role. It becomes necessary to document the algorithm, the implementation choices, and the range of operation. The general practice of "writing code that does not need inline documentation" does not apply. To an expert in the field, who has comprehensive understanding of the underlying math, such a code might be accessible without inline documentation. But to all others, and there are many who have reasons to look at the code, it would be equivalent to having equations without accompanying explanation. For longevity and extensibility, a scientific code must have inline documentation explaining the implementation logic, and reasons behind the choices made.

1.5   Case Study: The FLASH Code

1.5.1   Code Design

From the outset FLASH was required to have composability because the simulations of interest needed capabilities in different permutations and combinations. For example, most simulations needed compressible hydrodynamics, but with different equations of state. Some needed to include self-gravity while others did not. An obvious solution was to use object-oriented programming model with common API's and specializations to account for the different models. However, the physics capabilities were mostly legacy with F77 implementations. Rewriting the code in an object oriented language was not

an option. A compromise was found by exploiting the unix directory structure for inheritance, where, for a code unit the top level directory defined the API and the subdirectories contained the multiple alternative implementations of the API. Meta-information about the role of a particular directory level in the object oriented framework was encoded in a very limited domain-specific language (configuration DSL). The meta-information also included state and runtime variables requirements, dependences on other code units etc. A "setup tool" parsed this information to configure a consistent "application". The setup tool also interpreted the configuration DSL to implement inheritance using the directory structure. For more details about FLASH's object oriented framework see [**?**, **?**].

FLASH design is aware of the need for separation of concerns and achieves it by separating the infrastructural components from physics. The abstraction that permits this approach is very well known in CSE, that of decomposing a physical domain into rectangular blocks surrounded by halo cells copied over from the surrounding neighboring blocks. To a physics operator whole domain is not distinguishable from a box. Another necessary aspect of the abstraction is not to let any of the physics own the state variables. They are owned by the infrastructure that decomposes the domain into blocks. A further separation of concern takes place within the units handling the infrastructure, that of isolating parallelism from the bulk of the code. Parallel operations such as ghost cell fill, refluxing or regridding have minimal interleaving with state update in the blocks from application of physics operators. To distance the solvers from their parallel constructs, the required parallel operations provide an API with corresponding functions implemented as a subunit. The implementation of numerical algorithms for physics operators is sequential, interspersed with access to the parallel API as needed.

Minimization of data movement is achieved by letting the state be completely owned by the infrastructure modules. The dominant infrastructure module is the *Eulerian* mesh, owned and managed by the *Grid* unit. The physics modules query the *Grid* unit for the bounds and extent of the block they are operating on, and get a pointer to the physical data. This arrangement works in most cases, but gets tricky where the data access pattern does not conform to the underlying mesh. An example is any physics dealing with Lagrangian entities (LE's). They need a different data structure, and the movement of data has nothing in common with the way the data moves on the mesh. The added difficulty is that the entities do need to interact with the mesh, so physical proximity of the corresponding mesh cell is important in distributing the LE's. This is one of the examples of unavoidable lateral interaction between modules. In order to advance, LE's need to get some field quantities from the mesh and then determine their new locations internally. In some applications they have to apply near- and far-field forces, and in some applications they have to pass some information along to the mesh. And after advancing in time they may need to be redistributed. FLASH solves this conundrum through keeping the LE data structure extremely simple, and using

argument passing by reference in the API's. The LE's are attached to the block in the mesh that has the overlapping cell, an LE leaves its block when its location no longer overlaps with the block. Migration to a new block is an independent operation from everything else that goes on with the LE's In FLASH parlance this is the Lagrangian framework (see [**?**] for more details). The combination of *Eulerian* and *Lagrangian* frameworks that interoperate well with one another has succeeded in largely meeting the performance critical data management needs of the code.

1.5.2   Verification & Validation

FLASH instituted a rigorous verification program early in its lifecycle. The earliest versions of FLASH were subjected to a battery of standard hydrody-namics verification tests [**?**]. These verification tests were then used to set up an automated regression test suite run on a few local workstations. Since then the test suite has evolved into a combination of variety of tests that aim to provide comprehensive coverage for verifying correct operation of the code [**?, ?**]. Because FLASH is in a constant state of production and development, verification of its correctness on a regular basis is a critical requirement. The testing is complicated both by the variety of environments in which FLASH is run, and the sheer number of ways in which it can be configured.

FLASH's testing can be broadly classified into three categories: the daily testing to verify ongoing correctness of the code, more targeted testing related to science production runs, and finally porting to and testing on new platforms. Daily testing is performed on multiple combinations of platforms and software stacks. It includes unit tests where possible, and a modified version of no-change tests elsewhere. Because of floating point related issues mentioned in section **??** a drift within a specified tolerance is accepted as *no-change*. These tests also include *restart* tests, because a typical simulation is larger than one instance of batch queue allocation. It is imperative that the simulation be able to resume transparently from a checkpoint. All no-change tests for FLASH also incorporate verification of transparent restart for the corresponding problem.

FLASH selects problem setups for no-change testing using the matrix ap-proach to provide full block coverage [**?**]. The process of selecting problem setups is manual and relies upon the combined knowledge and expertise of code developers and in-house users. Following order is used for populating the matrix:

1. All unit tests.

2. Where possible problem setups corresponding to ongoing research sim-ulations.

3. Tests known to be sensitive to perturbations.

4. Tests known to exercise solvers in unusual ways.

5. Simplest problem setups that can fill the remaining gaps

In preparing for a production schedule, testing is a combination of scaling tests, cost estimation tests, and looking for potential trouble spots. Scientists and developers work closely to devise meaningful weak scaling tests (which can be difficult because of non-linearity and adaptive mesh refinement), and tests that can exercise the vulnerable code sections without overwhelming the test suite resources. Sample smaller scale production runs are also done on the target platform for make informed estimates of cpu hours and disk space needed to complete the simulation. For more details on simulation planning see [**?**]. For porting the code to a new platform a successful production run from the past is used as a benchmark for exercising the code on a new platform, along with using a subset of the standard test suite.

FLASH has had some opportunities for validation against experiments. For example FLASH could model a variety of laboratory experiments involving fluid instabilities [**?**, **?**]. These efforts allowed researchers to probe the validity of models and code modules, and also served to bolster the experimental efforts by creating realistic simulation capabilities for use in experimental design. The newer high-energy density physics (HEDP) initiative involving FLASH is directed at simulation-based validation and design of experiments at the major laser facilities in the US and Europe. Other forms of validation have been convergence tests for the flame model that is used for supernova simulations, and validation of various numerical algorithms against analytical solutions of some known problems. For example, the Sedov [] problem, which is seeded by a pressure spike in the center that sends out a spherical shock-wave into the domain has a known analytical solution. It is used to validate hydrodynamics in the code. There are several other similar examples where a simple problem can help to validate a code capability through known analytical solutions.

1.5.3   Software Process

The software process of FLASH has evolved organically with the growth of the code. For instance, in the first version there was no clear design document, the second version had a loosely implied design guidance, whereas the third version documented the whole design process. The third version also published the developer's guide which is a straight adaptation from the design document. Because of multiple developers with different production targets, versioning repository was introduced early in the code life cycle. The repository used has been SVN since 2003, though its branching system has been used in some very unorthodox ways to meet peculiar needs of the Flash Center. Unlike most software projects where branches are kept for somewhat isolated development purposes, FLASH uses branches also to manage multiple ongoing projects. This particular need arose when there were four different streams of physics capabilities being added to the code. All projects needed some code from the trunk, but the code being added was mostly exclusive to the individual project. It was important that the branches stay more or less in sync with the

trunk and that the new code be tested regularly. This was accomplished by turning the trunk into essentially a merge area, with a schedule of merge from individual branches, and an intermediate branch for forward merge. The path was tagged-trunk => forward-branch => projects => merge into trunk => tag trunk when stabilized. Note that the forward branch was never allowed a backward merge to avoid the possible inadvertent breaking of code for one project by another one. For the same reason the project branches never did a forward merge directly from the trunk.

One of the biggest challenges in managing a code like FLASH occurs during major version changes, when the infrastructure of the code undergoes deep changes. FLASH has undergone two such changes where the first transition took the approach of keeping the development branch synchronized with the main branch at all times. The new version also tried to keep itself backward compatible with the old version. During and after the process the team realized many shortcomings of this approach. One was that the code needed to have deeper structural changes than were possible under this approach. Also, the attempt to keep the development and production versions in sync placed undue demands on the developers of the new version, leading to inefficient use of their time. The adoption of the new version was delayed because keeping up with the ongoing modifications to the older version (needed by the scientists to do their work) turned the completion of the transition into a moving target.

Because of these lessons learned the second transition took a completely different tack and was much more successful. The infrastructural backbone/framework for the new version were built in isolation from the code base in a new repository. The framework design leveraged the knowledge gained by the developers about the idiosyncracies of the solvers in earlier versions and focussed on the needs of the future version. There was no attempt at backward compatibility with the framework of the previous version. Once the framework was thoroughly tested, physics modules were transitioned. Here the emphasis was on transitioning all the capabilities needed for one project at the same time, starting with the most stable modules. Once a module was moved to the new version it was effectively frozen in the old version (the reason for selecting the most stable and mature code sections). Any modification after that point had to be made simultaneously in the new version as well. Though it sounds like a lot of duplicate effort, in reality such instances were rare. This version transition was adopted by the scientists very quickly.

Testing is another area where the standard practices do not adequately meet the needs of the code. Many multiphysics codes have legacy components in them that are written in early versions of Fortran. Contrary to popular belief, a great deal of new development continues in Fortran because it still is the best HPC language in which to express mathematical algorithms. All of solver code in FLASH is F90, so the general unit test harnesses aren't available for use. Small scale unit tests can only be devised for infrastructural code because all the physics has to interact with the mesh. Also, because regular testing became a part of FLASH development process long before

formal incorporation of software engineering practices in the process, FLASH's designation of tests only loosely follows the standard definitions. So a unit test in FLASH can rely on other parts of the code, as long as the feature being tested is isolated. For example testing for correct filling of halo cells uses a lot of AMR code that has little to do with the halo filling, but it is termed unit test in FLASH parlance because it exclusively tests a single limited functionality. The dominant form of regular testing is integration testing, where more than one code capability is combined to configure an executable. The results of the run are compared against pre-approved results to verify that changes are within a specified acceptable range. Because of a large space of possible valid and useful combinations selection of tests is challenging. FLASH's methodology for test design and selection is described in detail in [**?**], and follows the matrix method described in chapter **??**

FLASH's documentation takes a comprehensive approach with a user's guide, a developer's guide, robodoc API, inline documentation, and online resources. Each type of documentation serves a different purpose and is indispensable to the developers and users of the code. There are scripts in place that look for violations of coding standards and documentation requirements. User's guide documents the mathematical formulation, algorithms used and instructions on using various code components. The user's guide also includes examples of relevant applications explaining the use of each code module. The developer's guide specifies the design principles and coding standards with an extensive example of the module architecture. Each function in the API is required to have a robodoc header explaining the input/output, function and special features of the function. Except for the third party software, every non-trivial function in the code is required to have sufficient inline documentation that a non-expert can understand and maintain the code.

FLASH effectively has two versions of release - internal, which is close to the agile model, and general, which is no more than twice a year. The internal release amounts to tagging a stable version in the repository for the internal users of the code. This is signal to the users that a forward merge into their production branch is not going to break the code. The general releases have a more rigorous process which makes them infrequent. The general releases undergo some amount of code pruning, get checked for compliance with coding and documentation standards and meet stringent requirements from the testing process. They are expensive in terms of developers resources. The dual model ensures that the quality of code and documentation are maintained without unduly straining the team resources, while near continuous code improvement is still possible for ongoing projects.

1.5.4   Policies

In any project, policies regarding attributions, contributions and licensing matter. In CSE arena intellectual property rights, and interdisciplinary interactions are additional policy areas that are equally important. Some of

these policy requirements are a direct consequence of the cathedral model [] of development that majority of CSE publicly distributed software follow. Many arguments are forwarded for dominance of the cathedral model in this domain, the most compelling one relates to maintaining the quality of software. Recollect that the developers in this domain are typically not trained in software engineering, and software quality control varies greatly between individuals and/or groups of developers. Because of tight, and sometimes lateral, coupling between functionalities of code modules a lower quality component introduced into the code base can have disproportionate impact on the overall reliability of output produced by the code. Strong gate-keeping is desirable, and that implies having policies in place for accepting contributions. FLASH again differentiates between internal and external contributors in this regard. The internal contributors are required to meet the quality requirements such as coding standards, documentation, and code verification in all of their development. Internal audit processes minimize the possibility of poorly written and tested code from getting into a release. The internal audit also goes through a periodic pruning to ensure that bad or redundant code gets eliminated.

The external contributors are required to work with a member of the internal team to include their code in the released version. The minimum set required from them is: (1) code that meets coding standards, has been used or will be used for results reported in peer-reviewed publication, (2) at least one test that can be included in the test-suite for nightly testing, (3) documentation for user's guide, robodoc documentation for any API functions and inline documentation explaining the flow of the control, and finally (4) a commitment to answer questions on users mailing list. The contributors can negotiate the terms of release, a code section can be excluded from the release for a mutually agreed period of time to enable the contributor to complete their research and publish their work before the code becomes public. This policy permits the potential contributors to be freed from the necessity of maintaining their code independently, while still retaining control over their software until agreed upon release time. As a useful side effect their code remains in sync with the developments in the main branch between releases.

There is another model of external contribution to FLASH that is without any intervention from the core gate-keeping team. In this model anyone can stage any FLASH compatible code on a site hosted by them. The code has no endorsement from the distributing entity, the Flash Center, which does not take any responsibility for its quality. The Flash Center maintains a list of externally hosted "as-is" code sites, the support for these code sections are entirely the responsibility of hosting site.

The attribution practices in CSE are somewhat ad-hoc. For many developers, the only metric of importance are the scientific publications that result from using the software. When a team is dominated by such developers proper attribution is not given enough importance or thought. Other teams also employ computing professionals whose career growth depends upon their software artifacts, and publications describing their algorithms and artifacts. FLASH

falls into the latter category, but the attribution policy does not reflect meet this challenge adequately. All contributors' names are included in the author list for the user's guide, the release notes explicitly mention new external contributions and their contributors, if any, for that release. Internal contributors rely upon software related publications for their attribution. This policy has not always worked well, and one of worst side effects has been citations skewed in favor of early developers. Users of FLASH cite a paper published in 2000 [**?**] which does not include any of the later code contributors in its author list, who are, therefore, deprived of legitimate recognition for citations. Many major long running software projects have this problem, which is peculiar to the academic world where these codes reside and are used.

1.6   Generalization

:

Not all of the solutions described in the earlier sections for CSE specific challenges are generalizable to all scientific software, but the vast majority of them are. This is borne out by the fact that at a workshop on community codes in 2012 [], all represented codes had nearly identical stories to tell about their motivation for adopting software engineering practices and the ones that they adopted. This was true irrespective of the science domains these codes served, the algorithms and discretization methods they used and communities they represented. Even their driving design principles were similar at the fundamental level though the details differed. The codes represented state-of-the-art in their respective communities in terms of both model and algorithmic research incorporated and the software engineering practices. Note that these are the codes that have stood the test of time and won the respect in their respective communities. They are widely used and supported, and have more credibility for producing reproducible reliable results than smaller individualistic efforts. Therefore, it is worthwhile to discuss those practices in this chapter. At a minimum they provide a snapshot of the state of large scale computing and its dependence of software engineering in the era of relatively uniform computing platforms.

One practice that is universally adopted by all community codes and other large scale codes is versioning repositories. That is worthy of mention here because even this practice has not penetrated the whole computational science community. There are many small projects that still do not use versioning, though their number is steadily decreasing. Other common practice is that of licensing for public use and most codes are freely available to download along with their source. Testing is also universal, though the extent and methodologies for testing vary greatly. A general verification and validation regime is still relatively rare, though regression testing is more common. Unit tests

are less common than integration tests and bounded-change tests. Almost all codes have user level documentation and user support practices in place. They also have well defined code contribution policies.

Another feature that stands out is the broader design philisophy of all multiphysics codes. Every code exercises separation of concerns between mathematical and structural parts and between sequential and parallel parts. In almost all cases this separation is dictated by the need to reduce complexity for efforts needing specific expertise. Also, all the codes have basic backbone frameworks which orchestrate the data movement and ownership. This is usually driven by the need for maintenance and flexibility. And where it is realized well it provides extensibility - the ability to add more physics and therefore greater capabilities and fidelity in the models being computed. Majority of frameworks are component based with composability of some sort. This is because different models need different capability combinations. Most codes use self-describing IO libraries for their output to facilitate the use of generally available analysis and visualization tools.

The degree to which teams from vastly different scientific domains producing community codes have arrived at essentially similar solutions is remarkable. It points to a possibility that seemingly diverse problems can have a uniform solution if they are trying to achieve similar objectives. For the codes highlighted in this section, the objectives were capabilities, extensibility, composability, reliability, portability and maintainability. They were achieved through design choices concious of trade-offs, most often with raw performance that individual components or specific platforms were capable of. The lesson here is that similar objectives can yield a general solution even if there is great diversity in the details of the individual problem. It is not beyond the realm of possibility that similar generalized solution will emerge for the next generation software faced with heterogeneous computing described in the next section.

1.7   Additional Future Considerations

:

One of the aspect of software design that is a unique requirement of the CSE domain is fast becoming its biggest challenge - performance portability. In the past machine architectures were fairly uniform across the board for large stretches of time. The first set of effective HPC machines in routine use for scientific computing were all vectors machines. They later gave way to parallel machines with *risc* processor as their main processing element. A code written for one machine of its time, if portable, would have reasonable performance on most of its contemporary machines. The abstract machine model to which the codes of the era were programming was essentially the

same for all machines of that era. It is true that wholesale changes had to occur in codes for transitioning from vector to risc-parallel machines, but it was a transition from one long-term stable paradigm to another long-term stable paradigm. And the codes were not as large as the multiphysics codes of today. So although the transitions took time, and the codes that adapted well to the prevailing machine model thrived for several years.

That landscape is about to change completely. Now there are machines in the pipeline that have deep enough architectural differences among them that one machine model cannot describe their behavior. Even within a machine heterogeneity of various kinds may exist. Because the codes are significantly larger than the last time such drastic changes had occurred in the computing platforms, the challenge is of a completely different magnitude. More importantly, some aspects of the challenges are not unique to the large multiphysics codes. Because the deep architectural changes are occurring at the level of nodes that will go into all platforms, the change is ubiquitous and will affect everyone. Portability in general and performance portability in particular is an issue for everyone. At this writing the impact of this paradigm shift is not fully understood. Means of combating this challenge are understoon even less. There is a general consensus that more programming abstractions are necessary not just for the extreme scale, but also for small scale computing. The unknown is which abstraction or combination of abstractions will deliver the solution. Many solutions have been proposed, for example [] (also see [**?**] for a more comprehensive and updated list). Of these, some have undergone more testing and exercise under realistic application instances than others. None of the approaches provide a good road map for a general solution that can be broadly applicable in the ways that optimizing compilers and MPI were in the past. This is an urgent serious challenge facing the CSE community today, future viability of CSE codes depends upon significant help from software engineering expertise and motivation within the community.