
Contents

1	Software Process for Multiphysics Multicomponent Codes	1
	<i>Anshu Dubey, Katherine Riley, and Katie Antypas</i>	
1.1	Introduction	1
1.2	Lifecycle	2
1.2.1	Development Cycle	2
1.2.2	Verification and Validation	3
1.2.2.1	Testing	4
1.2.3	Maintenance and Extensions	5
1.2.4	Performance Portability	5
1.2.5	Using CSE Software	6
1.3	Domain Challenges	6
1.4	Institutional Challenges	8
1.5	Case Study: The FLASH Code	10
1.5.1	Code Design	10
1.5.2	Verification & Validation	12
1.5.3	Software Process	13
1.5.4	Policies	15
1.6	Generalization	17
1.7	Additional Future Considerations	18
	Bibliography	21



Chapter 1

Software Process for Multiphysics Multicomponent Codes

1.1 Introduction

Computational science and engineering (CSE) communities develop complex applications to solve scientific and engineering challenges, but these communities have a mixed record of using software engineering best practices. Many codes developed by CSE communities adopt standard software practices when the size and complexity of an application becomes too unwieldy to continue without them [13]. The driving force behind adoption is usually the realization that without using software engineering practices, the development, the verification and the maintenance of applications can become intractable. As more codes cross the threshold into increasing complexity, software engineering processes are being adopted from practices derived outside of the scientific and engineering domain. State-of-the-art for software engineering practices in CSE codes often lags behind that in the commercial software space. There are many reasons for it, first, is that not all software engineering practices can actually be adopted by the CSE code developers. Secondly, there is often little funding for code development and maintenance. And finally, there is very little research in the software engineering community targeted to specific needs of CSE codes.

Many software best practices are not well-suited for CSE codes without modification and/or customization. In particular, multiphysics and multicomponents codes that run on the large High Performance Computing (HPC) platforms have specific requirements. In some cases, the inherent physics of scientific applications require different software methodologies. In other cases, a large premium is placed on performance, rather than code architecture, making it necessary to sacrifice some known software engineering best practices. Still others are more sociological because codes may be developed by domain scientists and their graduate students who have different priorities and scientific goals. The challenges in developing scientific and engineering applications range from data dependencies, to architectural trade-offs, to the process for their maintenance and growth. The standard practices adopted by the CSE codes include repositories for code version control, modular code design, licens-

ing process, regular testing, documentation, release and distribution policies, and contribution policies. Less frequently used practices include code-review, code-deprecation, and adoption of development methodologies such as Agile [1]. The degree of adoption and sophistication in using software engineering practices varies among teams. Many of the reasons for lower penetration of more formal software engineering practices are discussed in section 1.4. Even among the widely adopted practices, most are modified and customized by the developers for their own needs. The next few sections outline the challenges that are either unique to, or are more dominant in this domain than elsewhere.

1.2 Lifecycle

Scientific software is designed to model some phenomena in the physical world. The term 'physical' used here includes chemical and biological systems since physical processes are underlying building blocks for those systems too. A phenomenon may be at a microscopic level, for example protein folding, or at extremely large scales, for example galaxy cluster mergers. In some applications multiple scales are modeled. The physical characteristics of the systems being modeled are translated into mathematical models that are said to describe the essential features of the behavior of the system being studied. These equations are then discretized, and numerical algorithms are used to solve them. This process requires diverse expertise and adds many stages in the development and lifecycle of scientific software that may not be encountered elsewhere.

1.2.1 Development Cycle

For scientific simulations, modeling begins with equations that describe the general class of behavior to be studied, for example, the Navier-Stokes equations describe the flow of compressible and incompressible fluids, and van der Waals equations describe interactions among molecules in a material. There may be more than one set of equations if all behaviors of interest are not adequately captured by one set. In translating the model from mathematical representation to computational representation two processes go on simultaneously, discretization and approximation. One can argue that discretization is, by definition, an approximation because it is in effect sampling continuous behavior where information is lost in the sampling interval. This loss manifests itself as error terms in the discretized equations, but error terms are not the only approximations. Depending upon the level of understanding of specific sub-phenomena, and available compute resources, scientists also use their judgement to make other approximations. Sometimes, to focus on a particular behavior, a term in an equation may be simplified or may be even completely

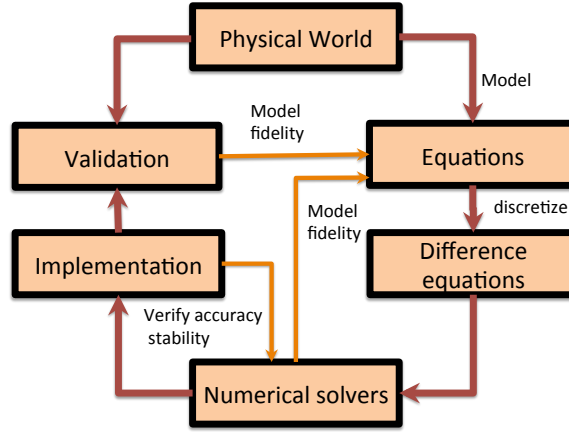


FIGURE 1.1: Development cycle of modeling with partial differential equations

dropped. At other times some physical details may be dropped from the model because they are not understood well enough by the scientists. Or the model itself may be an approximation.

The next stage in developing the code is finding appropriate numerical methods for each of the models. Sometimes existing methods can be used without modification, but more often, customization or new method development is needed. A method's applicability to the model may need to be validated if the method is new or significantly modified. Unless a reliable implementation is already available as a third party software (stand-alone or in a library), the method has to be implemented and verified. It is at this stage that the development of a CSE code begins to resemble that of general software. The numerical algorithms are specified, the semantics are understood, and they need to be translated into executable code. Figure 1.1 gives an example of the development cycle of a multiphysics application modeled using partial differential equations.

1.2.2 Verification and Validation

The terms verification and validation are often used interchangeably, but to some communities they have specific meaning. In one narrow definition, validation ensures that the mathematical model correctly defines the physical phenomena, while verification makes sure that the implementation of the model is correct. In other words, a model is validated against observations or experiments from the physical world, whereas a model is verified by other

forms of testing. Other definitions give broader scope to validation. For example, validation of a numerical method may be constructed through code-to-code comparisons, and its order can be validated through convergence studies. Similarly, the implementation of a solver can be validated against an analytically obtained solution for some model if the same solver can be applied and the analytical solution is also known, though this is not always possible. Irrespective of specific definitions, what is true is that correctness must be assured at all the stages from model to implementation.

There are many degrees of freedom in the process of deriving a model as discussed in the previous section, therefore, the validation of the model must be carefully calibrated by scientific experts. Similarly, verification of a numerical method requires applied math expertise, because, in addition to correctness the method needs verification of its stability, accuracy and order of convergence. Many numerical methods are themselves objects of ongoing research, so their implementation may need modifications from time to time. Whenever this happens, the entire gamut of verification and validation needs to be applied again. This is an instance of a particular challenge in CSE software where no amount of specification is enough to hand the implementation over to software engineers or developers who do not have domain or method knowledge. A close collaboration with applied mathematicians and method developers is necessary because the process has to be iterative with scientific judgement applied at every iteration.

One other unique verification challenge in CSE software is because of finite machine precision of floating point numbers. Any change in compilers, optimization levels, and even order of operations can cause numerical drift in the solutions. Especially in applications that have a large range of scales, it can be extremely difficult to differentiate between a legitimate bug and a numerical drift. Therefore, relying upon bitwise reproducibility of the solution is rarely a sufficient method for verifying the continued correctness of an application behavior. Robust diagnostics (such as statistics or conservation of physical quantities) need to be built into the verification process. This issue is discussed in greater detail in the chapter *Testing in CSE Software*.

1.2.2.1 Testing

Testing of CSE software needs to reflect the layered complexity that the codes themselves have. The first line of attack is to develop the unit tests which isolate testing of one component of the code. However, in CSE codes, often there are dependencies between different components of the code that can not be meaningfully isolated, making unit testing more difficult. In these cases, testing should be done with a minimal possible combination of components. In effect, these minimally combined tests play the same role in the testing regime that unit tests do because they focus on possible defects in a very narrow section of the code. In addition, multicomponent CSE software should test various permutations and combination of components in different ways.

Configuring tests in this manner will help verify that all configurations are within the accuracy and stability constraints.

1.2.3 Maintenance and Extensions

In a simplified view of software lifecycle, there is a design and development phase, followed by production and maintenance phase. Even in well engineered codes this simplified view typically applies only to the infrastructure and API's which have a distinct development phase with limited spill into the remainder of the lifecycle. The numerical algorithms and solvers can be in a continually evolving state reflecting the advances in their respective fields. The development of CSE software is usually responding to an immediate scientific need, so the codes get employed in production as soon as a minimal set of computational modules necessary for even one scientific project are built. Similarly, the development of computational modules almost never stops all through the code lifecycle because new findings in science and math almost continuously place new demands on them. The additions are mostly incremental when they incorporate new findings into an existing feature, they can also be substantial when new capabilities are added. The need for new capabilities may arise from greater model fidelity, or from trying to simulate a more complex model. Sometimes a code designed for one scientific field may have enough in common with another field that capabilities may be added to enable it for the new field.

Whatever may be the cause, co-existence of development and production/maintenance phases is a constant challenge to the code teams. It becomes acute when the code needs to undergo major version changes. The former can be managed with some repository discipline in the team coupled with a solid testing regime. The latter is a much bigger challenge where the plan has to concern itself with questions such as how much backward compatibility is suitable, how much code can go offline, and how to reconcile ongoing development in code sections that are substantially different between versions. FLASH's example in section 1.5.3 describes a couple of strategies that met the conflicting needs of developers and production users in both scenarios. Both required co-operation and buy-in from all the stakeholders to be successful.

1.2.4 Performance Portability

Another aspect of multiphysics CSE software is its need for performance portability. HPC machines are expensive and rare resources and in order to achieve high application performance, codes need to be optimized for the unique HPC architectures. However, typical lifecycle of a multiphysics application spans many generations of HPC systems which have a typical lifespan of about 4-5 years. Depending upon the size of the code, optimization for a specific target platform can take a significant fraction of the platform lifecycle, time when the code may not be available for science runs. Without careful

planning and coordination, a large fraction of scientists' time could be lost in porting and optimizing a code for a new system. Developers have the choice of adding machine specific optimizations or creating more general optimization that will work for a broader class of systems. CSE codes must consider the trade-offs and advantages of a highly optimized code for a single platform versus designing their software using constructs that perform modestly well across a range of platforms.

1.2.5 Using CSE Software

There is a fundamental requirement from the users of scientific software that rarely comes into play for users of other kinds of software. For obtaining valid scientific results the users must have a basic understanding of the phenomena being modeled and the approximations, and they must know which questions can be validly addressed by these models. For example, Eulerian equations of hydrodynamics do not account for viscosity. If one is modeling a fluid where viscosity is important, one should not use a code that uses Eulerian equations to model the fluid. Similarly, they must know and understand the valid regimes for applicability of numerical algorithms, and also the accuracy and stability behavior of the algorithms. For example a numerical method that can resolve a smooth fluid flow only will fail if there are discontinuities. Similarly, in order to use the Fast Fourier Transform (FFT) method, user must know the highest mode being resolved. If enough terms are not used in the Fourier series approximation the computed result may filter out important information and lead to wrong results.

Similarly, sometimes equations have mathematically valid but physically invalid solutions. An inappropriately applied numerical scheme may converge to such a non-physical solution. At best any of the above situations may lead to a waste of computing resources if the defect in the solution is detected. At worst they may lead to wrong scientific conclusions being drawn. Some in the scientific community even argue that those who have not written at least some basic version of their own code for their own problem should not be using other's public or community code. Though that argument goes too far, it embodies a general belief in the scientific community that users of scientific codes have a great responsibility to know their tools and understand their capabilities and limitations. These are some of the reasons that also play a role in tendency of scientific codes to do strict gatekeeping for contributions.

1.3 Domain Challenges

Multiphysics codes, by their definition, have more than one mathematical model that they are solving. A typical code combines 3-4 diverse models, the

more extreme ones may employ many more. In a rare calculation all models work with the same discretization using similar algorithmic approach (for instance stencil computations in explicit PDE solvers). More common is to have models with diverse discretizations and algorithms. In this case, each operator has its own preferred data layout and movement, and it usually differs from those needed by the other operators. Normally these challenges can be mitigated through encapsulation and well defined API's. The outer wrapper layers of the operators can carry out data transformations as needed. There are two factors against taking this approach in CSE codes: (1) physics is not always friendly to encapsulation, and (2) the codes are performance sensitive and wholesale data movement significantly degrades performance.

The CSE simulation codes model the physical world which does not have neat modularization. Various phenomena have tightly coupled dependencies that are hard to break. These dependencies and tight couplings also translate into their mathematical models and it becomes hard to eliminate lateral interactions among code modules implementing the models. An attempt to force encapsulations by hoisting up the lateral interactions to the API level can explode the size of the API. And if not done carefully this can also lead to extra data movement. The module designs, therefore, have to be cognizant of potential lateral interactions and make allowances for them. Similarly, the data structures have to take into account the diverse demands placed on them by different operators and carefully consider the trade-offs during software design. Considerations such as these are not common in software outside of CSE.

When designing CSE software, it is especially important to create modular components wherever possible. This is because different expertise may be required to understand each component, and thus a modular design allows application developers to focus on the areas they know best. For example, the numerical algorithms associated with physics operators require mathematical expertise, which is different from the code architecture which requires software engineering expertise. In addition, modular code allows various components to interface with each other in a clearer way. Another challenge for CSE codes is that many of them run on large HPC systems and require code parallelization. Parallel codes typically must implement a parallel domain decomposition and manage synchronization and load balancing between tasks or threads. With appropriate code modularization, these different aspects of the code do not need to interfere with one another and can help make application development more tractable.

Multiphysics multiscale codes often require tight integration with third party software, which comes in the form of numerical libraries. Because multiphysics codes combine expertise from many domains, the numerical solvers they use also require diverse applied mathematics expertise. It can be challenging for any one team to assemble all the necessary expertise to develop their own software and so many turn to third party math libraries for highly optimized routines. However, as mentioned in section 1.2.5, the use of third party software does not absolve them from understanding its appropriate use.

Additionally, information about appropriate use of third party software within the context of a larger code must also be communicated to the users of the code.

1.4 Institutional Challenges

Many adaptations in the software engineering for CSE applications described in the previous section pertained to software design and testing. In particular they spoke to challenges of modularity, performance and unit-testing because of the intertwined nature of the problems being tackled by these codes. However, there are a number of challenges that arise in developing and maintaining CSE codes outside of architectural and performance challenges. Many challenges are specific to the kind of organizations and the research communities where these codes are developed. The most crippling and pervasive challenge faced by CSE codes in general, and multiphysics codes in particular, is that funding for software development and maintenance is difficult to attain. There is evidence that when software is designed well it pays huge dividends in scientific productivity from the small number of projects that secured such funding for software infrastructure design. Examples include community codes such as NAMD [18], Amber [8] and Enzo [21] which are used by significant number of users in their respective communities. More persuasive case can be made by a handful of codes such as FLASH [10, 14], Cactus [6] and Uintah [17, 5] that were built for one community, but have since expanded their capabilities to serve several other communities using a common infrastructure. Even with this evidence it remains difficult to obtain funding for investment in software engineering best practices. Available funding is most often carved out of scientific goal oriented projects that have their own priorities and time-line. This model often ends up short-changing the software engineering.

The scientific output of applications is measured in terms of publications, which in turn depend upon the data produced by the simulations. Therefore, in a project driven purely by scientific objectives, the short-term science goals can lead to situations where quick-and-dirty triumphs over long term planning and design. The cost of future lost productivity may not be appreciated until much later when code base has grown too large to remove the deficiencies in any easy way. Software engineering is forcibly imposed on the code, which it at best a band-aid solution. This is another reason why many of the software practices are not embraced by the CSE community.

Another institutional challenge in developing good software engineering practices for CSE codes is training students and staff to use the application properly. Multiphysics codes require a broad range of expertise in domain science from their developers, and software engineering skills is an added requirement. Often experts in a domain science who develop CSE codes are not

trained in software engineering and many learn skills on the job through reading, or talking to colleagues. Practices are applied as they understand them, usually picking only what is of most importance for their own development. A very good example is the use of repositories. The “xkcd”¹ cartoon is essentially the truth of how git is used by many scientific software developers. This can be both good and bad. Good because it sifts out the unnecessary aspects of the SE practice, and bad because it is not always true that the sifted out aspects were really not necessary. It just means that the person adopting the practice did not understand the usefulness and impact of those aspects.

Institutional challenges also arise from scarcity and stability of resources. The domain and numerical algorithmic expertise is rarely replicated in a team developing the multiphysics CSE application. Even otherwise, deep expertise in the domain may be needed to model the phenomenon right, and that kind of expertise is hard to come by. Then there is the challenge of communicating the model to the software engineer, if there is one on the team, or to other members of the team with some other domain expertise. It requires at least a few developers in the team who can act as interpreters for various domain expertise and are able to integrate them. Such abilities take a great deal of time and effort to develop, neither of which are possible in academic institutions where these codes are typically organically grown. The available human resources in these institutions are post-docs and students who move on, so there is no retention of institutional knowledge about the code. A few projects that do see the need for software professionals struggle to find ways of funding them or to provide a path to professional growth.

The above institutional challenges also provide a clue about why any set software development methodology is hard, and often even undesirable, to adopt in such projects. For example, the principles behind the agile manifesto apply, but not all the formalized processes do. Agile software methods [1] are lightweight evolutionary development methods where focus is on adaptability and flexibility, as opposed to waterfall methods which is a sequential development process where progress is perceived as a downward flow [3]. Agile methods aim to deliver working software as early as possible within the lifecycle and improve it based upon user feedback and changing needs. These aims fit well with the objectives of scientific software development as well. These codes are developed by interdisciplinary teams where interactions and collaborations are preferred over regimented process. The code development and its use for science go on in parallel, so the requirements change and there is quick feedback when they do. For the same reason, the code needs to be in working condition almost all the time. However, scarcity of resources does not allow the professional roles in the agile process to be played out efficiently. There is no clear separation between the developer and the client, many developers of the code are also the scientists who use it for their research. Because software development goes hand in hand with research and exploration of the

¹<https://xkcd.com/1597/>

algorithms it is impossible to do either within fixed time-frames. This constraint effectively eliminates using agile methods such as *sprints* [1]. Similarly, extreme programming, another component of Agile, is impossible to implement in an environment which has to incorporate research into the process. The waterfall model is even less useful because it is impossible to do a full specification ahead of time. The code has to grow and alter organically as the scientific understanding grows, the effect of using technologies are digested and requirements change.

The need for deep expertise, and the fact that the developer of a complex physics module is almost definitely going to leave with possibly no replacement, documentation of various kind takes on a crucial role. It becomes necessary to document the algorithm, the implementation choices, and the range of operation. The generally preferred practice of writing self explanatory code helps, but does not suffice. To an expert in the field, who has comprehensive understanding of the underlying math, such a code might be accessible without inline documentation. But not to non-experts (i.e. from another field or a software engineer in the team if there is one) who may have reasons to look at the code. For longevity and extensibility, a scientific code must have inline documentation explaining the implementation logic, and reasons behind the choices made.

1.5 Case Study: The FLASH Code

1.5.1 Code Design

From the outset FLASH was required to have composability because the simulations of interest needed capabilities in different permutations and combinations. For example, most simulations needed compressible hydrodynamics, but with different equations of state. Some needed to include self-gravity while others did not. An obvious solution was to use object-oriented programming model with common API's and specializations to account for the different models. However, the physics capabilities were mostly legacy with F77 implementations. Rewriting the code in an object oriented language was not an option. A compromise was found by exploiting the unix directory structure for inheritance, where, for a code unit the top level directory defined the API and the subdirectories contained the multiple alternative implementations of the API. Meta-information about the role of a particular directory level in the object oriented framework was encoded in a very limited domain-specific language (configuration DSL). The meta-information also included state and runtime variables requirements, dependences on other code units etc. A "setup tool" parsed this information to configure a consistent "application". The setup tool also interpreted the configuration DSL to implement

inheritance using the directory structure. For more details about FLASH's object oriented framework see [10].

FLASH design is aware of the need for separation of concerns and achieves it by separating the infrastructural components from physics. The abstraction that permits this approach is very well known in CSE, that of decomposing a physical domain into rectangular blocks surrounded by halo cells copied over from the surrounding neighboring blocks. To a physics operator whole domain is not distinguishable from a box. Another necessary aspect of the abstraction is not to let any of the physics own the state variables. They are owned by the infrastructure that decomposes the domain into blocks. A further separation of concern takes place within the units handling the infrastructure, that of isolating parallelism from the bulk of the code. Parallel operations such as ghost cell fill, refluxing or regridding have minimal interleaving with state update in the blocks from application of physics operators. To distance the solvers from their parallel constructs, the required parallel operations provide an API with corresponding functions implemented as a subunit. The implementation of numerical algorithms for physics operators is sequential, interspersed with access to the parallel API as needed.

Minimization of data movement is achieved by letting the state be completely owned by the infrastructure modules. The dominant infrastructure module is the *Eulerian* mesh, owned and managed by the *Grid* unit. The physics modules query the *Grid* unit for the bounds and extent of the block they are operating on, and get a pointer to the physical data. This arrangement works in most cases, but gets tricky where the data access pattern does not conform to the underlying mesh. An example is any physics dealing with Lagrangian entities (LE's). They need a different data structure, and the movement of data has nothing in common with the way the data moves on the mesh. The added difficulty is that the entities do need to interact with the mesh, so physical proximity of the corresponding mesh cell is important in distributing the LE's. This is one of the examples of unavoidable lateral interaction between modules. In order to advance, LE's need to get some field quantities from the mesh and then determine their new locations internally. In some applications they have to apply near- and far-field forces, and in some applications they have to pass some information along to the mesh. And after advancing in time they may need to be redistributed. FLASH solves this conundrum through keeping the LE data structure extremely simple, and using argument passing by reference in the API's. The LE's are attached to the block in the mesh that has the overlapping cell, an LE leaves its block when its location no longer overlaps with the block. Migration to a new block is an independent operation from everything else that goes on with the LE's. In FLASH parlance this is the Lagrangian framework (see [12] for more details). The combination of *Eulerian* and *Lagrangian* frameworks that interoperate well with one another has succeeded in largely meeting the performance critical data management needs of the code.

1.5.2 Verification & Validation

FLASH instituted a rigorous verification program early in its lifecycle. The earliest versions of FLASH were subjected to a battery of standard hydrodynamics verification tests [15]. These verification tests were then used to set up an automated regression test suite run on a few local workstations. Since then the test suite has evolved into a combination of variety of tests that aim to provide comprehensive coverage for verifying correct operation of the code [11, 7]. Because FLASH is in a constant state of production and development, verification of its correctness on a regular basis is a critical requirement. The testing is complicated both by the variety of environments in which FLASH is run, and the sheer number of ways in which it can be configured.

FLASH's testing can be broadly classified into three categories: the daily testing to verify ongoing correctness of the code, more targeted testing related to science production runs, and finally porting to and testing on new platforms. Daily testing is performed on multiple combinations of platforms and software stacks. It includes unit tests where possible, and a modified version of no-change tests elsewhere. Because of floating point related issues mentioned in section 1.2 a drift within a specified tolerance is accepted as *no-change*. These tests also include *restart* tests, because a typical simulation is larger than one instance of batch queue allocation. It is imperative that the simulation be able to resume transparently from a checkpoint. All no-change tests for FLASH also incorporate verification of transparent restart for the corresponding problem.

FLASH selects problem setups for no-change testing using the matrix approach to provide full block coverage [19]. The process of selecting problem setups is manual and relies upon the combined knowledge and expertise of code developers and in-house users. Following order is used for populating the matrix:

1. All unit tests.
2. Where possible problem setups corresponding to ongoing research simulations.
3. Tests known to be sensitive to perturbations.
4. Tests known to exercise solvers in unusual ways.
5. Simplest problem setups that can fill the remaining gaps

In preparing for a production schedule, testing is a combination of scaling tests, cost estimation tests, and looking for potential trouble spots. Scientists and developers work closely to devise meaningful weak scaling tests (which can be difficult because of non-linearity and adaptive mesh refinement), and tests that can exercise the vulnerable code sections without overwhelming the test suite resources. Sample smaller scale production runs are also done on the target platform for make informed estimates of cpu hours and disk space needed to complete the simulation. For more details on simulation planning see

[11]. For porting the code to a new platform a successful production run from the past is used as a benchmark for exercising the code on a new platform, along with using a subset of the standard test suite.

FLASH has had some opportunities for validation against experiments. For example FLASH could model a variety of laboratory experiments involving fluid instabilities [9, 16]. These efforts allowed researchers to probe the validity of models and code modules, and also served to bolster the experimental efforts by creating realistic simulation capabilities for use in experimental design. The newer high-energy density physics (HEDP) initiative involving FLASH is directed at simulation-based validation and design of experiments at the major laser facilities in the US and Europe. Other forms of validation have been convergence tests for the flame model that is used for supernova simulations, and validation of various numerical algorithms against analytical solutions of some known problems. For example, the Sedov [20] problem, which is seeded by a pressure spike in the center that sends out a spherical shock-wave into the domain has a known analytical solution. It is used to validate hydrodynamics in the code. There are several other similar examples where a simple problem can help to validate a code capability through known analytical solutions.

1.5.3 Software Process

The software process of FLASH has evolved organically with the growth of the code. For instance, in the first version there was no clear design document, the second version had a loosely implied design guidance, whereas the third version documented the whole design process. The third version also published the developer's guide which is a straight adaptation from the design document. Because of multiple developers with different production targets, versioning repository was introduced early in the code life cycle. The repository used has been SVN since 2003, though its branching system has been used in some very unorthodox ways to meet peculiar needs of the Flash Center. Unlike most software projects where branches are kept for somewhat isolated development purposes, FLASH uses branches also to manage multiple ongoing projects. This particular need arose when there were four different streams of physics capabilities being added to the code. All projects needed some code from the trunk, but the code being added was mostly exclusive to the individual project. It was important that the branches stay more or less in sync with the trunk and that the new code be tested regularly. This was accomplished by turning the trunk into essentially a merge area, with a schedule of merge from individual branches, and an intermediate branch for forward merge. The path was tagged-trunk => forward-branch => projects => merge into trunk => tag trunk when stabilized. Note that the forward branch was never allowed a backward merge to avoid the possible inadvertent breaking of code for one project by another one. For the same reason the project branches never did a forward merge directly from the trunk.

One of the biggest challenges in managing a code like FLASH occurs dur-

ing major version changes, when the infrastructure of the code undergoes deep changes. FLASH has undergone two such changes where the first transition took the approach of keeping the development branch synchronized with the main branch at all times. The new version also tried to keep itself backward compatible with the old version. During and after the process the team realized many shortcomings of this approach. One was that the code needed to have deeper structural changes than were possible under this approach. Also, the attempt to keep the development and production versions in sync placed undue demands on the developers of the new version, leading to inefficient use of their time. The adoption of the new version was delayed because keeping up with the ongoing modifications to the older version (needed by the scientists to do their work) turned the completion of the transition into a moving target.

Because of these lessons learned the second transition took a completely different tack and was much more successful. The infrastructural backbone/framework for the new version were built in isolation from the code base in a new repository. The framework design leveraged the knowledge gained by the developers about the idiosyncracies of the solvers in earlier versions and focussed on the needs of the future version. There was no attempt at backward compatibility with the framework of the previous version. Once the framework was thoroughly tested, physics modules were transitioned. Here the emphasis was on transitioning all the capabilities needed for one project at the same time, starting with the most stable modules. Once a module was moved to the new version it was effectively frozen in the old version (the reason for selecting the most stable and mature code sections). Any modification after that point had to be made simultaneously in the new version as well. Though it sounds like a lot of duplicate effort, in reality such instances were rare. This version transition was adopted by the scientists very quickly.

Testing is another area where the standard practices do not adequately meet the needs of the code. Many multiphysics codes have legacy components in them that are written in early versions of Fortran. Contrary to popular belief, a great deal of new development continues in Fortran because it still is the best HPC language in which to express mathematical algorithms. All of solver code in FLASH is F90, so the general unit test harnesses aren't available for use. Small scale unit tests can only be devised for infrastructural code because all the physics has to interact with the mesh. Also, because regular testing became a part of FLASH development process long before formal incorporation of software engineering practices in the process, FLASH's designation of tests only loosely follows the standard definitions. So a unit test in FLASH can rely on other parts of the code, as long as the feature being tested is isolated. For example testing for correct filling of halo cells uses a lot of AMR code that has little to do with the halo filling, but it is termed unit test in FLASH parlance because it exclusively tests a single limited functionality. The dominant form of regular testing is integration testing, where more than one code capability is combined to configure an executable. The results of the run are compared against pre-approved results to verify that changes are within a

specified acceptable range. Because of a large space of possible valid and useful combinations selection of tests is challenging. FLASH's methodology for test design and selection is described in detail in [4], and follows the matrix method described earlier.

FLASH's documentation takes a comprehensive approach with a user's guide, a developer's guide, robodoc API, inline documentation, and online resources. Each type of documentation serves a different purpose and is indispensable to the developers and users of the code. There are scripts in place that look for violations of coding standards and documentation requirements. User's guide documents the mathematical formulation, algorithms used and instructions on using various code components. The user's guide also includes examples of relevant applications explaining the use of each code module. The developer's guide specifies the design principles and coding standards with an extensive example of the module architecture. Each function in the API is required to have a robodoc header explaining the input/output, function and special features of the function. Except for the third party software, every non-trivial function in the code is required to have sufficient inline documentation that a non-expert can understand and maintain the code.

FLASH effectively has two versions of release - internal, which is close to the agile model, and general, which is no more than twice a year. The internal release amounts to tagging a stable version in the repository for the internal users of the code. This is signal to the users that a forward merge into their production branch is not going to break the code. The general releases have a more rigorous process which makes them infrequent. The general releases undergo some amount of code pruning, get checked for compliance with coding and documentation standards and meet stringent requirements from the testing process. They are expensive in terms of developers resources. The dual model ensures that the quality of code and documentation are maintained without unduly straining the team resources, while near continuous code improvement is still possible for ongoing projects.

1.5.4 Policies

In any project, policies regarding attributions, contributions and licensing matter. In CSE arena intellectual property rights, and interdisciplinary interactions are additional policy areas that are equally important. Some of these policy requirements are a direct consequence of the strong gatekeeping regimes that majority of CSE publicly distributed software follow. Many arguments are forwarded for dominance of this model in the domain, the most compelling one relates to maintaining the quality of software. Recollect that the developers in this domain are typically not trained in software engineering, and software quality control varies greatly between individuals and/or groups of developers. Because of tight, and sometimes lateral, coupling between functionalities of code modules a lower quality component introduced into the code base can have disproportionate impact on the overall reliability of output produced by

the code. Strong gate-keeping is desirable, and that implies having policies in place for accepting contributions. FLASH again differentiates between internal and external contributors in this regard. The internal contributors are required to meet the quality requirements such as coding standards, documentation, and code verification in all of their development. Internal audit processes minimize the possibility of poorly written and tested code from getting into a release. The internal audit also goes through a periodic pruning to ensure that bad or redundant code gets eliminated.

The external contributors are required to work with a member of the internal team to include their code in the released version. The minimum set required from them is: (1) code that meets coding standards, has been used or will be used for results reported in peer-reviewed publication, (2) at least one test that can be included in the test-suite for nightly testing, (3) documentation for user's guide, robodoc documentation for any API functions and inline documentation explaining the flow of the control, and finally (4) a commitment to answer questions on users mailing list. The contributors can negotiate the terms of release, a code section can be excluded from the release for a mutually agreed period of time to enable the contributor to complete their research and publish their work before the code becomes public. This policy permits the potential contributors to be freed from the necessity of maintaining their code independently, while still retaining control over their software until agreed upon release time. As a useful side effect their code remains in sync with the developments in the main branch between releases.

There is another model of external contribution to FLASH that is without any intervention from the core gate-keeping team. In this model anyone can stage any FLASH compatible code on a site hosted by them. The code has no endorsement from the distributing entity, the Flash Center, which does not take any responsibility for its quality. The Flash Center maintains a list of externally hosted 'as-is' code sites, the support for these code sections are entirely the responsibility of hosting site.

The attribution practices in CSE are somewhat ad-hoc. For many developers, the only metric of importance are the scientific publications that result from using the software. When a team is dominated by such developers proper attribution for code development is not given enough importance or thought. Other teams also employ computing professionals whose career growth depends upon their software artifacts, and publications describing their algorithms and artifacts, FLASH falls into the latter category. All contributors' names are included in the author list for the user's guide, the release notes explicitly mention new external contributions and their contributors, if any, for that release. Internal contributors rely upon software related publications for their attribution. This policy usually works well, one exception has been citations skewed in favor of early developers. Users of FLASH cite a paper published in 2000 [15] which does not include any of the later code contributors in its author list. So even though the current incarnation of FLASH has very little of the original code, the original authors continue to get the

benefit while the later authors are deprived of legitimate citations for their work. Many major long running software projects have this problem, which is peculiar to the academic world where these codes reside and are used.

1.6 Generalization

:

Not all of the solutions described in the earlier sections for CSE specific challenges are generalizable to all scientific software, but the vast majority of them are. This is borne out by the fact that at a workshop on community codes in 2012 [13], all represented codes had nearly identical stories to tell about their motivation for adopting software engineering practices and the ones that they adopted. This was true irrespective of the science domains these codes served, the algorithms and discretization methods they used and communities they represented. Even their driving design principles were similar at the fundamental level though the details differed. The codes represented state-of-the-art in their respective communities in terms of both model and algorithmic research incorporated and the software engineering practices. Note that these are the codes that have stood the test of time and won the respect in their respective communities. They are widely used and supported, and have more credibility for producing reproducible reliable results than smaller individualistic efforts. Therefore, it is worthwhile to discuss those practices in this chapter. At a minimum they provide a snapshot of the state of large scale computing and its dependence on software engineering in the era of relatively uniform computing platforms.

One practice that is universally adopted by all community codes and other large scale codes is versioning repositories. That is worthy of mention here because even this practice has not penetrated the whole computational science community. There are many small projects that still do not use versioning, though their number is steadily decreasing. Other common practice is that of licensing for public use and most codes are freely available to download along with their source. Testing is also universal, though the extent and methodologies for testing vary greatly. A general verification and validation regime is still relatively rare, though regression testing is more common. Unit tests are less common than integration tests and bounded-change tests. Almost all codes have user level documentation and user support practices in place. They also have well defined code contribution policies.

Another feature that stands out is the broader design philosophy of all multiphysics codes. Every code exercises separation of concerns between mathematical and structural parts and between sequential and parallel parts. In almost all cases this separation is dictated by the need to reduce complexity for efforts needing specific expertise. Also, all the codes have basic backbone

frameworks which orchestrate the data movement and ownership. This is usually driven by the need for maintenance and flexibility. And where it is realized well it provides extensibility - the ability to add more physics and therefore greater capabilities and fidelity in the models being computed. Majority of frameworks are component based with composability of some sort. This is because different models need different capability combinations. Most codes use self-describing IO libraries for their output to facilitate the use of generally available analysis and visualization tools.

The degree to which teams from vastly different scientific domains producing community codes have arrived at essentially similar solutions is remarkable. It points to a possibility that seemingly diverse problems can have a uniform solution if they are trying to achieve similar objectives. For the codes highlighted in this section, the objectives were capabilities, extensibility, composability, reliability, portability and maintainability. They were achieved through design choices conscious of trade-offs, most often with raw performance that individual components or specific platforms were capable of. The lesson here is that similar objectives can yield a general solution even if there is great diversity in the details of the individual problem. It is not beyond the realm of possibility that similar generalized solution will emerge for the next generation software faced with heterogeneous computing described in the next section.

1.7 Additional Future Considerations

:
One of the aspect of software design that is a unique requirement of the CSE domain is fast becoming its biggest challenge - performance portability. In the past, machine architectures were fairly uniform across the board for large stretches of time. The first set of effective HPC machines in routine use for scientific computing were all vectors machines. They later gave way to parallel machines with *risc* processor as their main processing element. A code written for one machine of its time, if portable, would have reasonable performance on most of its contemporary machines. The abstract machine model to which the codes of the era were programming was essentially the same for all machines of that era. It is true that wholesale changes had to occur in codes for transitioning from vector to risc-parallel machines, but it was a transition from one long-term stable paradigm to another long-term stable paradigm. And the codes were not as large as the multiphysics codes of today. So although the transitions took time, the codes that adapted well to the prevailing machine model thrived for several years.

That landscape is about to change completely. Now there are machines in the pipeline that have deep enough architectural differences among them that

one machine model cannot describe their behavior. Even within a machine heterogeneity of various kinds may exist. Because the codes are significantly larger than they were during the last architecture paradigm shift, the challenge is of a completely different magnitude. More importantly, some aspects of the challenges are not unique to the large multiphysics codes. Because the deep architectural changes are occurring at the level of nodes that will go into all platforms, the change is ubiquitous and will affect everyone. Portability in general and performance portability in particular is an issue for everyone. At this writing the impact of this paradigm shift is not fully understood. Means of combating this challenge are understood even less. There is a general consensus that more programming abstractions are necessary not just for the extreme scale, but also for small scale computing. The unknown is which abstraction or combination of abstractions will deliver the solution. Many solutions have been proposed, for example [22] (also see [2] for a more comprehensive and updated list). Of these, some have undergone more testing and exercise under realistic application instances than others. None of the approaches provide a good road map for a general solution that can be broadly applicable in the ways that optimizing compilers and MPI were in the past. This is an urgent serious challenge facing the CSE community today, future viability of CSE codes depends upon significant help from software engineering expertise and motivation within the community.



Bibliography

- [1] Agile methodology. <http://agilemethodology.org/>.
- [2] Ideas productivity - howto documents. <https://ideas-productivity.org/resources/howtos/>.
- [3] Waterfall model. <https://www.techopedia.com/definition/14025/waterfall-model>.
- [4] A. Dubey, K. Weide, D. Lee, J. Bachan, C. Daley, S. Olofin, N. Taylor, P.M. Rich, and L.B. Reid. Ongoing verification of a multiphysics community code: FLASH. *Software: Practice and Experience*, 45(2), 2015.
- [5] M. Berzins, J. Luitjens, Q. Meng, T. Harman, C.A. Wight, and J.R. Peterson. Uintah - a scalable framework for hazard analysis. In *TG '10: Proc. of 2010 TeraGrid Conference*, New York, NY, USA, 2010. ACM.
- [6] Marek Blazewicz, Ian Hinder, David M Koppelman, Steven R Brandt, Milosz Ciznicki, Michal Kierzynka, Frank Löffler, Erik Schnetter, and Jian Tao. From physics model to results: An optimizing framework for cross-architecture code generation. *Scientific Programming*.
- [7] A. C. Calder. Laboratory astrophysics experiments for simulation code validation: A case study. *Astrophysics and Space Science*, 298:25–32, July 2005.
- [8] DA Case, V Babin, Josh Berryman, RM Betz, Q Cai, DS Cerutti, TE Cheatham Iii, TA Darden, RE Duke, H Gohlke, et al. Amber 14. 2014.
- [9] G. Dimonte, D. L. Youngs, A. Dimits, S. Weber, M. Marinak, S. Wunsch, C. Garasi, A. Robinson, M. J. Andrews, P. Ramaprabhu, A. C. Calder, B. Fryxell, J. Biello, L. Dursi, P. MacNeice, K. Olson, P. Ricker, R. Rosner, F. Timmes, H. Tufo, Y.-N. Young, and M. Zingale. A comparative study of the turbulent Rayleigh-Taylor instability using high-resolution three-dimensional numerical simulations: The Alpha-Group collaboration. *Physics of Fluids*, 16:1668–1693, May 2004.
- [10] A. Dubey, K. Antypas, M.K. Ganapathy, L.B. Reid, K. Riley, D. Sheeler, A. Siegel, and K. Weide. Extensible component-based architecture for

- FLASH, a massively parallel, multiphysics simulation code. *Parallel Computing*, 35(10-11):512–522, 2009.
- [11] A. Dubey, A.C. Calder, C. Daley, R.T. Fisher, C. Graziani, G.C. Jordan, D.Q. Lamb, L.B. Reid, D. M. Townsley, and K. Weide. Pragmatic optimizations for better scientific utilization of large supercomputers. *International Journal of High Performance Computing Applications*, 27(3):360–373, 2013.
 - [12] A. Dubey, C. Daley, J. ZuHone, P. M. Ricker, K. Weide, and C. Graziani. Imposing a Lagrangian particle framework on an Eulerian hydrodynamics infrastructure in FLASH. *ApJ Supplement*, 201:27, aug 2012.
 - [13] A. Dubey, D.Q. Lamb, and E. Balaras. Building community codes for effective scientific research on HPC platforms. <http://flash.uchicago.edu/cc2012>, 2012.
 - [14] A. Dubey, L.B. Reid, and R. Fisher. Introduction to FLASH 3.0, with application to supersonic turbulence. *Physica Scripta*, T132, 2008. Topical Issue on Turbulent Mixing and Beyond, results of a conference at ICTP, Trieste, Italy, August 2008.
 - [15] B. Fryxell, K. Olson, P. Ricker, F. X. Timmes, M. Zingale, D. Q. Lamb, P. MacNeice, R. Rosner, J. W. Truran, and H. Tufo. Flash: An adaptive mesh hydrodynamics code for modeling astrophysical thermonuclear flashes. *Astrophysical Journal, Supplement*, 131:273–334, 2000.
 - [16] JO Kane, HF Robey, BA Remington, RP Drake, J. Knauer, DD Ryutov, H. Louis, R. Teyssier, O. Hurricane, D. Arnett, et al. Interface imprinting by a rippled shock using an intense laser. *Physical review. E, Statistical, nonlinear, and soft matter physics*, 63(5 Pt 2):055401, 2001.
 - [17] S. G. Parker. A component-based architecture for parallel multi-physics PDE simulation. *Future Generation Comput. Sys.*, 22:204–216, 2006.
 - [18] James C Phillips, Rosemary Braun, Wei Wang, James Gumbart, Emad Tajkhorshid, Elizabeth Villa, Christophe Chipot, Robert D Skeel, Laxmikant Kale, and Klaus Schulten. Scalable molecular dynamics with namd. *Journal of computational chemistry*, 26(16):1781–1802, 2005.
 - [19] Xiao Qu, Myra B. Cohen, and Gregg Rothermel. Configuration-aware regression testing: an empirical study of sampling and prioritization. In *Proceedings of the 2008 international symposium on Software testing and analysis*, ISSTA '08, pages 75–86, New York, NY, USA, 2008. ACM.
 - [20] L.I. Sedov. Similarity and dimensional methods in mechanics.
 - [21] The Enzo Collaboration, G. L. Bryan, M. L. Norman, B. W. O’Shea, T. Abel, J. H. Wise, M. J. Turk, D. R. Reynolds, D. C. Collins, P. Wang,

- S. W. Skillman, B. Smith, R. P. Harkness, J. Bordner, J.-h. Kim, M. Kuhlen, H. Xu, N. Goldbaum, C. Hummels, A. G. Kritsuk, E. Tasker, S. Skory, C. M. Simpson, O. Hahn, J. S. Oishi, G. C So, F. Zhao, R. Cen, and Y. Li. Enzo: An Adaptive Mesh Refinement Code for Astrophysics. *ArXiv e-prints*, July 2013.
- [22] Didem Unat, J Shalf, T Hoefer, T Schulthess, A Dubey, et al. Programming abstractions for data locality. In *Workshop on programming abstractions for data locality (PADAL 14)*, 2014.